

# Assignment 2

Daniele Algeri, Martina Kenna, Joana Pimenta and Matilde Simonini

Master's Degree in Artificial Intelligence, University of Bologna

{ daniele.algeri, martina.kenna, joana.monteiro, matilde.simonini } @studio.unibo.it

## Abstract

This work compares two open-source LLMs, *TinyLlama-1.1B-Chat-v1.0* and *Mistral-7B-Instruct-v0.3*, on multiclass sexism detection using zero-shot and few-shot prompting. Results highlight a critical dependency on model size: TinyLlama consistently failed to follow instructions in both inference settings and, even after Prompt Tuning, semantic reasoning was still absent, leading to a single-class prediction bias, although formatting errors were corrected. In contrast, Mistral-7B demonstrated effective instruction adherence. Although initially biased towards one class, Mistral's performance improved significantly with few-shot examples, raising the F1-score from 0.35 to 0.46.

## 1 Introduction

Traditional approaches for text classification often rely on supervised learning with models like BERT (Jacob Devlin, Ming-Wei Chang and Kenton Lee, 2018), which require extensive fine-tuning. Recently, Large Language Models (LLMs) have introduced the paradigm of "prompting," allowing models to perform classification tasks without weight updates, relying instead on in-context learning. In this work, we investigate the efficacy of LLMs in categorizing a set of texts into five distinct categories: *not-sexist*, or one among four sexist classes, namely *threats*, *derogation*, *animosity*, and *prejudiced discussion*. We compare a lightweight model (TinyLlama-1.1B) against a larger model (Mistral-7B). Our experimental setup involves classifying a balanced, supervised test set of 300 texts, using two different strategies: **Zero-Shot** prompting, where the model is given the task description and input text only, and **Few-Shot** prompting, where the model is provided with the task description and a small set of labeled examples.

In addition, we perform **Prompt Tuning** to TinyLlama to address its instruction-following limitations. Our experiment highlights that:

- The smaller model (TinyLlama, 1.1B parameters) lacks the capacity to follow the syntactic constraints out-of-the-box, and outputs invalid labels for every instance in the dataset.
- Prompt Tuning successfully forced such model to output valid labels, but could not induce the reasoning required for correct classification.
- Mistral tends to default to the *animosity* class in the zero-shot setting.

## 2 System description

We implemented an inference pipeline based on the Hugging Face `transformers` library. The pipeline for each example in the dataset consists of the following stages:

1. **Prompt Construction:** We use a strict template defining the role of the model, listing the five valid categories, and instructing the model to output *only* the label.

- **Zero-Shot:** Instruction + target text.
- **Few-Shot:** Instruction + labeled examples + target text.

2. **Inference:** The prompt is passed to the LLM.

3. **Output Parsing:** The raw text response is mapped to numerical IDs (0–4). If the model generates invalid text, the system defaults to **0** (*not-sexist*).

To address TinyLlama limitations, we implemented an additional Prompt Tuning pipeline for such model. This technique involves freezing the pre-trained model backbone and optimizing only a small set of continuous vector embeddings that are prepended to the input (corresponding to approx. 0.006% of parameters), aiming to align the model with the specific prompt structure.

### 3 Experimental setup and results

The performance is evaluated on the Macro F1 Score and the Accuracy Score. The following tables summarize the performance metrics obtained with the two models.

Table 1: **Zero-Shot** Inference Metrics

<b>Model</b>	<b>F1-Score</b>	<b>Accuracy</b>
Mistral	0.35	0.38
TinyLlama	0.07	0.20

Table 2: **Few-Shot** Inference Metrics

<b>Model</b>	<b>F1-Score</b>	<b>Accuracy</b>
Mistral	0.46	0.48
TinyLlama	0.07	0.20

It is crucial to note that TinyLlama’s metrics are artifacts of the pipeline defaulting invalid outputs to class 0. While Prompt Tuning successfully corrected the syntax, it resulted in a semantic collapse where the model explicitly predicted *not-sexist* for nearly all inputs. Consequently, the metrics remained identical, effectively shifting from a fallback-driven default to a learned prediction bias without improving discrimination.

## 4 Discussion

Our results show that Mistral achieves a respectable zero-shot performance which is further boosted by few-shot examples (F1 increase from 0.35 to 0.46). TinyLlama, conversely, shows static metrics (Accuracy 0.20, F1 0.07) across all setups, which match the statistical probability of a random classifier for a balanced 5-class dataset.

### 4.1 Error Analysis

The invalid outputs provided by TinyLlama, in both Zero-Shot and Few-Shot Inference, consist in long generative text rather than discrete class labels.

A possible reason for this behavior is that TinyLlama-Chat is trained primarily for conversational settings, where verbose and explanatory answers are generally preferred. Even the inclusion of in-context demonstrations (Few-Shot prompting) proved insufficient to enforce syntactic adherence, likely because this technique relies on a model’s

<b>Text:</b> "damn, i bet she likes tall dudes"
<b>True Label:</b> "not-sexist"
<b>Predicted:</b> "Sexist Categories Definition: Threats: The text expresses intent [...]"

Figure 1: Example of invalid TinyLlama output

ability to perform in-context learning, which is known to scale strongly with model size.

Ultimately, the Prompt Tuning experiment was successful in making the model learn the syntax, as all 300 examples in the test dataset were classified with a single, valid label; nevertheless, the model still failed to learn the semantics, leading to unsatisfying performance once again. This dual failure—first in formatting, then in discrimination—strongly suggests that the 1.1B parameter size acts as a fundamental bottleneck for the classification performance.

Concerning Mistral, the model presents a strong bias towards the *animosity* class in the Zero-Shot setting, which, however, was significantly mitigated by few-shot prompting, as shown by the two confusion matrices in Figure 2

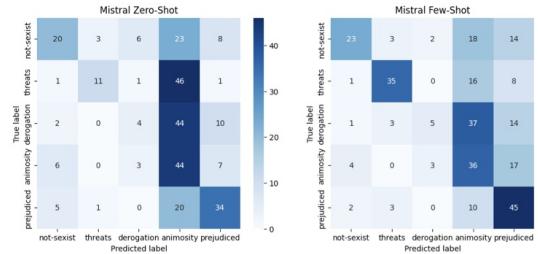


Figure 2: Confusion matrices for Mistral in Zero-Shot and Few-Shot Inference.

## 5 Conclusion

In this experiment, we evaluated the capabilities of LLMs for multiclass sexism detection. We observed that model size is a decisive factor not only for rigid instruction following (syntax) but also for the semantic understanding required to distinguish between complex classes. TinyLlama-1.1B proved unable to adhere to the output format in standard prompting; furthermore, even when Prompt Tuning successfully enforced syntactic compliance, the model lacked the representational depth to discriminate between categories, collapsing into a single-class prediction.

In contrast, Mistral-7B demonstrated that a suf-

ficiently large model can effectively interpret constraints and leverage context. The transition from zero-shot to few-shot proved highly effective for Mistral, rectifying a strong bias toward *animosity* and allowing for better identification of specific categories like *threats*. This confirms that larger models possess the necessary capacity to refine decision boundaries via in-context learning—a capability that appears absent in the smaller 1.1B architecture.

## References

- Jacob Devlin, Ming-Wei Chang and Kenton Lee. 2018.  
Bert: Pre-training of deep bidirectional transformers for language understanding. <https://arxiv.org/abs/1810.04805>.